

lecture06-penalized-regression-annotated.R

mac

2019-07-26

```
# STA303 S19 Lecture 6: penalized regression
#
# # Load required packages
#

library(dplyr)

##
## Attaching package: 'dplyr'

## The following objects are masked from 'package:stats':
##
##   filter, lag

## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union

library(ggplot2)
library(glmnet)

## Loading required package: Matrix
## Loading required package: foreach
## Loaded glmnet 2.0-18

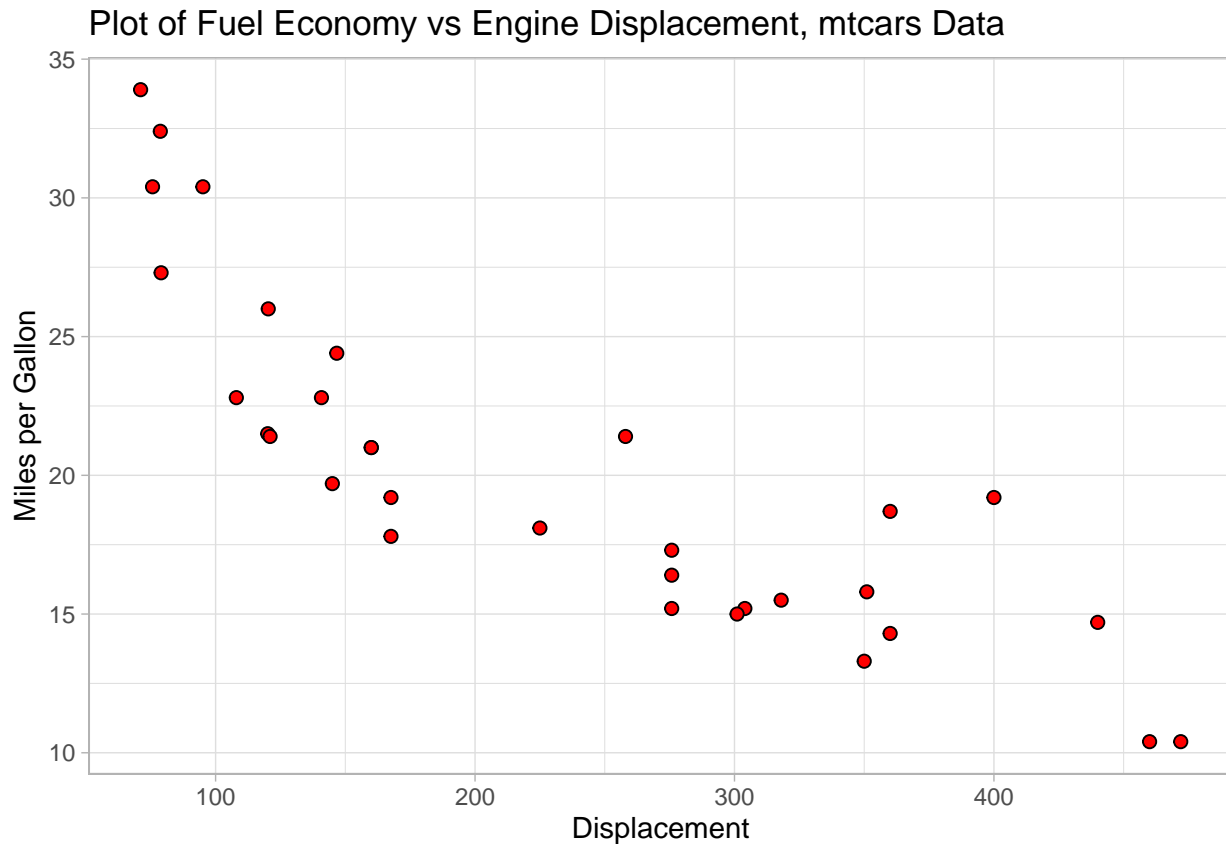
# Linear regression on the mtcars data.
mtcars <- mtcars %>%
  as_tibble() %>% # Removes rownames and pretty-prints
  select(displacement, mpg)
mtcars

## # A tibble: 32 x 2
##   displacement mpg
##   <dbl> <dbl>
## 1 160    21
## 2 160    21
## 3 108   22.8
## 4 258   21.4
## 5 360   18.7
## 6 225   18.1
## 7 360   14.3
## 8 147   24.4
## 9 141   22.8
## 10 168   19.2
## # ... with 22 more rows

# First thing to do: plot the data, see what type of relationship might be present.
# Plot
mt_plot <- mtcars %>%
```

```
ggplot(aes(x=disp,y=mpg)) +
  theme_light() +
  geom_point(colour="black",fill="red",pch=21,size=2) +
  labs(title = "Plot of Fuel Economy vs Engine Displacement, mtcars Data",
       x="Displacement",
       y="Miles per Gallon")
```

mt_plot



```
# Looks curvy. Linear regression?
# Scale the data first so that y and x have mean 0 and variance 1
# This is because the response and covariate are on different scales.
# So coefficient estimates are not immediately interpretable.
# Rescaling the data mitigates this.
```

```
mtscaled <- mtcars %>%
  mutate_all(~(.x - mean(.x))/sd(.x))
```

mtscaled

```
## # A tibble: 32 x 2
##       disp    mpg
##   <dbl> <dbl>
## 1 -0.571  0.151
## 2 -0.571  0.151
## 3 -0.990  0.450
## 4  0.220  0.217
## 5  1.04   -0.231
```


```
## 6 -0.0462 -0.330
## 7 1.04 -0.961
## 8 -0.678 0.715
## 9 -0.726 0.450
## 10 -0.509 -0.148
## # ... with 22 more rows

mod1 <- lm(mpg~disp,data=mtscaled)
summary(mod1)

##
## Call:
## lm(formula = mpg ~ disp, data = mtscaled)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.8117 -0.3654 -0.1598  0.2700  1.1997
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -7.037e-17  9.537e-02   0.000      1
## disp        -8.476e-01  9.689e-02  -8.747 9.38e-10 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.5395 on 30 degrees of freedom
## Multiple R-squared:  0.7183, Adjusted R-squared:  0.709
## F-statistic: 76.51 on 1 and 30 DF,  p-value: 9.38e-10

# We have centred the data, so an intercept of 0 means that a car with average
# displacement (engine size) has average mpg (fuel economy).
# Further, beta(disp) = -.848. A one standard deviation increase in disp is associated with a .847
# standard deviation decrease in fuel economy, on average.
# Predicted values:
mtscaled %>%
  ggplot(aes(x=disp,y=mpg)) +
  theme_light() +
  geom_point(colour="black",fill="red",pch=21,size=2) +
  geom_abline(intercept = coef(mod1)[1],slope=coef(mod1)[2],colour = "purple") +
  labs(subtitle="Linear Regression Model")
```


Linear Regression Model



*# Definitely looks like a linear relationship is missing some of the structure.
Can we pick this up in a residual plot?*

```
tibble(x = mod1$fitted.values,  
       y = residuals(mod1)) %>%  
  ggplot(aes(x=x,y=y)) +  
  theme_light() +  
  geom_point(colour="black",fill="red",pch=21,size=2) +  
  geom_hline(yintercept = 0,colour = "red") +  
  labs(title = "Residual plot, linear model",  
       x = "Fitted value",  
       y = "Residual")
```

Residual plot, linear model



```
# The scale of the residuals looks pretty good- all fall
# between -2 and 2, which we expect since the data
# was scaled to have unit variance.
# But the pattern? Doesn't look evenly distributed about
# y = 0. Looks like... a quadratic?
```

```
# Model fit- R2
summary(mod1)$r.squared # Fits pretty well.
```

```
## [1] 0.7183433
```

```
# Now try a quadratic...
# The poly() function just creates a polynomial out of
# the input variable with the specified degree.
# By default it creates something called "orthogonal polynomials",
# but you can get regular ones using raw = TRUE.
mod2 <- lm(mpg~poly(displacement,degree = 2,raw = TRUE),data=mtscaled)
summary(mod2)
```


```
##
## Call:
## lm(formula = mpg ~ poly(displacement, degree = 2, raw = TRUE), data = mtscaled)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.64895 -0.25334 -0.05184  0.22381  0.89509
##
## Coefficients:
```

```
##                                Estimate Std. Error t value Pr(>|t|)
## (Intercept)                   -0.30996    0.12710  -2.439   0.0211
## poly(displacement, degree = 2, raw = TRUE)1 -0.97360    0.09313 -10.454 2.39e-11
## poly(displacement, degree = 2, raw = TRUE)2  0.31996    0.09917   3.226   0.0031
##
## (Intercept)                    *
## poly(displacement, degree = 2, raw = TRUE)1 ***
## poly(displacement, degree = 2, raw = TRUE)2 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.4707 on 29 degrees of freedom
## Multiple R-squared:  0.7927, Adjusted R-squared:  0.7784
## F-statistic: 55.46 on 2 and 29 DF,  p-value: 1.229e-10

# Harder to directly interpret the betas in a polynomial model
# but the standard errors should be not huge compared to the point estimates
# if the fitting procedure was nice and stable.
# (remember: fitting  $y = \beta_0 + \beta_1 x + \beta_2 x^2$ ) is STILL linear regression
# because it's a linear function of the betas).

mtscaled %>%
  mutate(predvals = predict(mod2)) %>%
  ggplot(aes(x=displacement,y=mpg)) +
  theme_light() +
  geom_point(colour="black",fill="red",pch=21,size=2) +
  geom_line(aes(y = predvals),colour = "purple") +
  labs(subtitle="Linear Regression Model with quadratic term")
```


Linear Regression Model with quadratic term



*# A quadratic gives a much better fit, but it's not that flexible.
 # We see it starting to curve back up at the right-most x value
 # Would a more flexible (higher degree) polynomial be appropriate?*

```
# Residuals:
tibble(x = mod2$fitted.values,
       y = residuals(mod2)) %>%
  ggplot(aes(x=x,y=y)) +
  theme_light() +
  geom_point(colour="black",fill="red",pch=21,size=2) +
  geom_hline(yintercept = 0,colour = "red") +
  labs(title = "Residual plot, quadratic model",
       x = "Fitted value",
       y = "Residual")
```

Residual plot, quadratic model



I don't see anything that alarming in this plot. It looks good!

Coefficients

```
coef(mod2)
```

```
##               (Intercept) poly(dis, degree = 2, raw = TRUE)1
##               -0.3099566                -0.9736036
## poly(dis, degree = 2, raw = TRUE)2
##               0.3199552
```

Rsquared

```
summary(mod2)$r.squared # Fits a bit better than linear model.
```

```
## [1] 0.7927323
```

Adding a quadratic term improved the fit a bit.

Does adding more terms always improve the fit?

```
mod3 <- lm(mpg~poly(dis,degree = 20,raw = TRUE),data=mtscaled)
summary(mod3)
```

```
##
```

```
## Call:
```

```
## lm(formula = mpg ~ poly(dis, degree = 20, raw = TRUE), data = mtscaled)
```

```
##
```

```
## Residuals:
```

```
##      Min       1Q   Median       3Q      Max
```

```
## -0.4857 -0.1734  0.0000  0.1001  0.5695
```

```
##
```



```
## Coefficients: (2 not defined because of singularities)
##               Estimate Std. Error t value
## (Intercept)    -1.501e-02  3.904e-01  -0.038
## poly(dis, degree = 20, raw = TRUE)1    6.736e+00  4.874e+00   1.382
## poly(dis, degree = 20, raw = TRUE)2   -1.205e+01  2.765e+01  -0.436
## poly(dis, degree = 20, raw = TRUE)3   -9.245e+01  6.158e+01  -1.501
## poly(dis, degree = 20, raw = TRUE)4    6.054e+01  2.687e+02   0.225
## poly(dis, degree = 20, raw = TRUE)5    4.684e+02  4.223e+02   1.109
## poly(dis, degree = 20, raw = TRUE)6   -8.656e+01  1.071e+03  -0.081
## poly(dis, degree = 20, raw = TRUE)7   -1.262e+03  1.472e+03  -0.858
## poly(dis, degree = 20, raw = TRUE)8   -3.787e+01  2.229e+03  -0.017
## poly(dis, degree = 20, raw = TRUE)9    1.971e+03  2.780e+03   0.709
## poly(dis, degree = 20, raw = TRUE)10   2.029e+02  2.627e+03   0.077
## poly(dis, degree = 20, raw = TRUE)11  -1.843e+03  3.034e+03  -0.607
## poly(dis, degree = 20, raw = TRUE)12  -1.888e+02  1.762e+03  -0.107
## poly(dis, degree = 20, raw = TRUE)13   1.030e+03  1.950e+03   0.528
## poly(dis, degree = 20, raw = TRUE)14   7.138e+01  6.252e+02   0.114
## poly(dis, degree = 20, raw = TRUE)15  -3.302e+02  7.104e+02  -0.465
## poly(dis, degree = 20, raw = TRUE)16  -9.542e+00  9.093e+01  -0.105
## poly(dis, degree = 20, raw = TRUE)17   5.452e+01  1.276e+02   0.427
## poly(dis, degree = 20, raw = TRUE)18           NA           NA           NA
## poly(dis, degree = 20, raw = TRUE)19  -3.460e+00  7.245e+00  -0.478
## poly(dis, degree = 20, raw = TRUE)20           NA           NA           NA
##               Pr(>|t|)
## (Intercept)           0.970
## poly(dis, degree = 20, raw = TRUE)1    0.190
## poly(dis, degree = 20, raw = TRUE)2    0.670
## poly(dis, degree = 20, raw = TRUE)3    0.157
## poly(dis, degree = 20, raw = TRUE)4    0.825
## poly(dis, degree = 20, raw = TRUE)5    0.287
## poly(dis, degree = 20, raw = TRUE)6    0.937
## poly(dis, degree = 20, raw = TRUE)7    0.407
## poly(dis, degree = 20, raw = TRUE)8    0.987
## poly(dis, degree = 20, raw = TRUE)9    0.491
## poly(dis, degree = 20, raw = TRUE)10   0.940
## poly(dis, degree = 20, raw = TRUE)11   0.554
## poly(dis, degree = 20, raw = TRUE)12   0.916
## poly(dis, degree = 20, raw = TRUE)13   0.606
## poly(dis, degree = 20, raw = TRUE)14   0.911
## poly(dis, degree = 20, raw = TRUE)15   0.650
## poly(dis, degree = 20, raw = TRUE)16   0.918
## poly(dis, degree = 20, raw = TRUE)17   0.676
## poly(dis, degree = 20, raw = TRUE)18           NA
## poly(dis, degree = 20, raw = TRUE)19   0.641
## poly(dis, degree = 20, raw = TRUE)20           NA
##
## Residual standard error: 0.3544 on 13 degrees of freedom
## Multiple R-squared:  0.9473, Adjusted R-squared:  0.8744
## F-statistic: 12.99 on 18 and 13 DF, p-value: 1.496e-05
```


```
# Woah. The coefficients (point estimates) are all over the place in magnitude,
# and direction. Their standard errors are huge.
# This is the "instability" we talked about in lecture.
# Whether it's statistical (sampling) error or numerical error, something is wrong here.
```

```

# Rather than plot a linear/quadratic using geom_line(),
# to do predictions for 20 degree polynomial, I use the predict() function
# and then plot its output.
mtscaled %>%
  mutate(predvals = predict(mod3)) %>%
  ggplot(aes(x=disp,y=mpg)) +
  theme_light() +
  geom_point(colour="black",fill="red",pch=21,size=2) +
  geom_line(aes(y = predvals),colour = "purple") +
  labs(subtitle="Linear Regression Model with 20-degree polynomial terms")

```

Linear Regression Model with 20-degree polynomial terms



```

# While this model fits THESE data well, it fits them TOO well.
# Sacrifices generalizability to new datasets.
# Statistically: point estimates have very high variance. Confidence intervals
# would be wide. We're asking too much of this small dataset.
# Coefficients
unnname(coef(mod3))

```

```

## [1] -1.501106e-02  6.736184e+00 -1.204588e+01 -9.244740e+01  6.053962e+01
## [6]  4.683826e+02 -8.655610e+01 -1.262487e+03 -3.786733e+01  1.971390e+03
## [11]  2.028601e+02 -1.842881e+03 -1.887617e+02  1.029692e+03  7.137544e+01
## [16] -3.301832e+02 -9.542259e+00  5.451536e+01          NA -3.460233e+00
## [21]          NA

```

```

# Rsquared
summary(mod3)$r.squared # YES, this model does fit THESE DATA very well.

```

```
## [1] 0.9473424
```

```
# The model fits the observed data better. But it gives ridiculous predictions.
# The coefficients are huge. The fitted curve hugs the observed data too closely.
# This model will not generalize well to new data, and the coefficients are highly
# sensitive/unstable (look at their standard errors/p-values)
```

Penalized Regression

```
# glmnet() wants X and y to be provided separately (no formula)
# so get the design matrices from each of the above models.
# If this is unfamiliar, print out each of the below X's and see what's inside.
y <- mtscaled$mpg
y
```

```
## [1] 0.15088482 0.15088482 0.44954345 0.21725341 -0.23073453
## [6] -0.33028740 -0.96078893 0.71501778 0.44954345 -0.14777380
## [11] -0.38006384 -0.61235388 -0.46302456 -0.81145962 -1.60788262
## [16] -1.60788262 -0.89442035 2.04238943 1.71054652 2.29127162
## [21] 0.23384555 -0.76168319 -0.81145962 -1.12671039 -0.14777380
## [26] 1.19619000 0.98049211 1.71054652 -0.71190675 -0.06481307
## [31] -0.84464392 0.21725341
```

```
X1 <- model.matrix(mod1)
X1
```

```
##      (Intercept)      disp
## 1             1 -0.57061982
## 2             1 -0.57061982
## 3             1 -0.99018209
## 4             1  0.22009369
## 5             1  1.04308123
## 6             1 -0.04616698
## 7             1  1.04308123
## 8             1 -0.67793094
## 9             1 -0.72553512
## 10            1 -0.50929918
## 11            1 -0.50929918
## 12            1  0.36371309
## 13            1  0.36371309
## 14            1  0.36371309
## 15            1  1.94675381
## 16            1  1.84993175
## 17            1  1.68856165
## 18            1 -1.22658929
## 19            1 -1.25079481
## 20            1 -1.28790993
## 21            1 -0.89255318
## 22            1  0.70420401
## 23            1  0.59124494
## 24            1  0.96239618
## 25            1  1.36582144
## 26            1 -1.22416874
## 27            1 -0.89093948
## 28            1 -1.09426581
## 29            1  0.97046468
```

```
## 30      1 -0.69164740
## 31      1  0.56703942
## 32      1 -0.88529152
## attr(,"assign")
## [1] 0 1
```

```
X2 <- model.matrix(mod2)
X2
```

```
##      (Intercept) poly(displacement, degree = 2, raw = TRUE)1
## 1      1      -0.57061982
## 2      1      -0.57061982
## 3      1      -0.99018209
## 4      1       0.22009369
## 5      1       1.04308123
## 6      1      -0.04616698
## 7      1       1.04308123
## 8      1      -0.67793094
## 9      1      -0.72553512
## 10     1      -0.50929918
## 11     1      -0.50929918
## 12     1       0.36371309
## 13     1       0.36371309
## 14     1       0.36371309
## 15     1       1.94675381
## 16     1       1.84993175
## 17     1       1.68856165
## 18     1      -1.22658929
## 19     1      -1.25079481
## 20     1      -1.28790993
## 21     1      -0.89255318
## 22     1       0.70420401
## 23     1       0.59124494
## 24     1       0.96239618
## 25     1       1.36582144
## 26     1      -1.22416874
## 27     1      -0.89093948
## 28     1      -1.09426581
## 29     1       0.97046468
## 30     1      -0.69164740
## 31     1       0.56703942
## 32     1      -0.88529152
##      poly(displacement, degree = 2, raw = TRUE)2
## 1      0.32560698
## 2      0.32560698
## 3      0.98046057
## 4      0.04844123
## 5      1.08801845
## 6      0.00213139
## 7      1.08801845
## 8      0.45959036
## 9      0.52640121
## 10     0.25938565
## 11     0.25938565
## 12     0.13228721
```

```
## 13          0.13228721
## 14          0.13228721
## 15          3.78985041
## 16          3.42224749
## 17          2.85124044
## 18          1.50452130
## 19          1.56448766
## 20          1.65871200
## 21          0.79665117
## 22          0.49590329
## 23          0.34957057
## 24          0.92620640
## 25          1.86546820
## 26          1.49858911
## 27          0.79377315
## 28          1.19741766
## 29          0.94180170
## 30          0.47837612
## 31          0.32153370
## 32          0.78374108
## attr("assign")
## [1] 0 1 1
```

```
X3 <- model.matrix(mod3)
X3
```

```
##      (Intercept) poly(displacement, degree = 20, raw = TRUE)1
## 1             1          -0.57061982
## 2             1          -0.57061982
## 3             1          -0.99018209
## 4             1           0.22009369
## 5             1           1.04308123
## 6             1          -0.04616698
## 7             1           1.04308123
## 8             1          -0.67793094
## 9             1          -0.72553512
## 10            1          -0.50929918
## 11            1          -0.50929918
## 12            1           0.36371309
## 13            1           0.36371309
## 14            1           0.36371309
## 15            1           1.94675381
## 16            1           1.84993175
## 17            1           1.68856165
## 18            1          -1.22658929
## 19            1          -1.25079481
## 20            1          -1.28790993
## 21            1          -0.89255318
## 22            1           0.70420401
## 23            1           0.59124494
## 24            1           0.96239618
## 25            1           1.36582144
## 26            1          -1.22416874
## 27            1          -0.89093948
## 28            1          -1.09426581
```

```

## 29          1          0.97046468
## 30          1         -0.69164740
## 31          1          0.56703942
## 32          1         -0.88529152
##  poly(dis, degree = 20, raw = TRUE)2
## 1          0.32560698
## 2          0.32560698
## 3          0.98046057
## 4          0.04844123
## 5          1.08801845
## 6          0.00213139
## 7          1.08801845
## 8          0.45959036
## 9          0.52640121
## 10         0.25938565
## 11         0.25938565
## 12         0.13228721
## 13         0.13228721
## 14         0.13228721
## 15         3.78985041
## 16         3.42224749
## 17         2.85124044
## 18         1.50452130
## 19         1.56448766
## 20         1.65871200
## 21         0.79665117
## 22         0.49590329
## 23         0.34957057
## 24         0.92620640
## 25         1.86546820
## 26         1.49858911
## 27         0.79377315
## 28         1.19741766
## 29         0.94180170
## 30         0.47837612
## 31         0.32153370
## 32         0.78374108
##  poly(dis, degree = 20, raw = TRUE)3
## 1         -1.857978e-01
## 2         -1.857978e-01
## 3         -9.708345e-01
## 4          1.066161e-02
## 5          1.134892e+00
## 6         -9.839983e-05
## 7          1.134892e+00
## 8         -3.115705e-01
## 9         -3.819226e-01
## 10        -1.321049e-01
## 11        -1.321049e-01
## 12         4.811459e-02
## 13         4.811459e-02
## 14         4.811459e-02
## 15         7.377906e+00
## 16         6.330924e+00

```

```

## 17          4.814495e+00
## 18         -1.845430e+00
## 19         -1.956853e+00
## 20         -2.136272e+00
## 21         -7.110535e-01
## 22          3.492171e-01
## 23          2.066818e-01
## 24          8.913775e-01
## 25          2.547896e+00
## 26         -1.834526e+00
## 27         -7.072038e-01
## 28         -1.310293e+00
## 29          9.139853e-01
## 30         -3.308676e-01
## 31          1.823223e-01
## 32         -6.938393e-01
##   poly(dis, degree = 20, raw = TRUE)4
## 1          1.060199e-01
## 2          1.060199e-01
## 3          9.613029e-01
## 4          2.346553e-03
## 5          1.183784e+00
## 6          4.542823e-06
## 7          1.183784e+00
## 8          2.112233e-01
## 9          2.770982e-01
## 10         6.728092e-02
## 11         6.728092e-02
## 12         1.749991e-02
## 13         1.749991e-02
## 14         1.749991e-02
## 15         1.436297e+01
## 16         1.171178e+01
## 17         8.129572e+00
## 18         2.263584e+00
## 19         2.447622e+00
## 20         2.751325e+00
## 21         6.346531e-01
## 22         2.459201e-01
## 23         1.221996e-01
## 24         8.578583e-01
## 25         3.479972e+00
## 26         2.245769e+00
## 27         6.300758e-01
## 28         1.433809e+00
## 29         8.869904e-01
## 30         2.288437e-01
## 31         1.033839e-01
## 32         6.142501e-01
##   poly(dis, degree = 20, raw = TRUE)5
## 1         -6.049706e-02
## 2         -6.049706e-02
## 3         -9.518650e-01
## 4          5.164616e-04

```

```

## 5      1.234783e+00
## 6     -2.097284e-07
## 7      1.234783e+00
## 8     -1.431948e-01
## 9     -2.010445e-01
## 10    -3.426612e-02
## 11    -3.426612e-02
## 12     6.364945e-03
## 13     6.364945e-03
## 14     6.364945e-03
## 15     2.796116e+01
## 16     2.166599e+01
## 17     1.372728e+01
## 18    -2.776488e+00
## 19    -3.061472e+00
## 20    -3.543459e+00
## 21    -5.664616e-01
## 22     1.731779e-01
## 23     7.224989e-02
## 24     8.255995e-01
## 25     4.753020e+00
## 26    -2.749201e+00
## 27    -5.613594e-01
## 28    -1.568968e+00
## 29     8.607929e-01
## 30    -1.582792e-01
## 31     5.862276e-02
## 32    -5.437904e-01
##      poly(dis, degree = 20, raw = TRUE)6
## 1      3.452082e-02
## 2      3.452082e-02
## 3      9.425196e-01
## 4      1.136699e-04
## 5      1.287979e+00
## 6      9.682527e-09
## 7      1.287979e+00
## 8      9.707619e-02
## 9      1.458648e-01
## 10     1.745170e-02
## 11     1.745170e-02
## 12     2.315014e-03
## 13     2.315014e-03
## 14     2.315014e-03
## 15     5.443349e+01
## 16     4.008060e+01
## 17     2.317936e+01
## 18     3.405611e+00
## 19     3.829274e+00
## 20     4.563657e+00
## 21     5.055971e-01
## 22     1.219526e-01
## 23     4.271738e-02
## 24     7.945538e-01
## 25     6.491776e+00

```



```

## 26          3.365485e+00
## 27          5.001373e-01
## 28          1.716868e+00
## 29          8.353691e-01
## 30          1.094734e-01
## 31          3.324142e-02
## 32          4.814130e-01
##      poly(displacement, degree = 20, raw = TRUE)7
## 1          -1.969826e-02
## 2          -1.969826e-02
## 3          -9.332661e-01
## 4           2.501804e-05
## 5           1.343467e+00
## 6          -4.470130e-10
## 7           1.343467e+00
## 8          -6.581095e-02
## 9          -1.058301e-01
## 10         -8.888139e-03
## 11         -8.888139e-03
## 12          8.420008e-04
## 13          8.420008e-04
## 14          8.420008e-04
## 15          1.059686e+02
## 16          7.414638e+01
## 17          3.913979e+01
## 18         -4.177286e+00
## 19         -4.789636e+00
## 20         -5.877579e+00
## 21         -4.512723e-01
## 22          8.587949e-02
## 23          2.525643e-02
## 24          7.646756e-01
## 25          8.866607e+00
## 26         -4.119922e+00
## 27         -4.455920e-01
## 28         -1.878710e+00
## 29          8.106962e-01
## 30         -7.571697e-02
## 31          1.884919e-02
## 32         -4.261909e-01
##      poly(displacement, degree = 20, raw = TRUE)8
## 1           1.124022e-02
## 2           1.124022e-02
## 3           9.241033e-01
## 4           5.506312e-06
## 5           1.401345e+00
## 6           2.063724e-11
## 7           1.401345e+00
## 8           4.461528e-02
## 9           7.678343e-02
## 10          4.526722e-03
## 11          4.526722e-03
## 12          3.062467e-04
## 13          3.062467e-04

```

```

## 14 3.062467e-04
## 15 2.062948e+02
## 16 1.371657e+02
## 17 6.608994e+01
## 18 5.123814e+00
## 19 5.990852e+00
## 20 7.569792e+00
## 21 4.027846e-01
## 22 6.047668e-02
## 23 1.493274e-02
## 24 7.359209e-01
## 25 1.211020e+01
## 26 5.043480e+00
## 27 3.969955e-01
## 28 2.055808e+00
## 29 7.867520e-01
## 30 5.236945e-02
## 31 1.068824e-02
## 32 3.773032e-01
## poly(dis, degree = 20, raw = TRUE)9
## 1 -6.413892e-03
## 2 -6.413892e-03
## 3 -9.150306e-01
## 4 1.211905e-06
## 5 1.461717e+00
## 6 -9.527590e-13
## 7 1.461717e+00
## 8 -3.024608e-02
## 9 -5.570908e-02
## 10 -2.305456e-03
## 11 -2.305456e-03
## 12 1.113859e-04
## 13 1.113859e-04
## 14 1.113859e-04
## 15 4.016052e+02
## 16 2.537473e+02
## 17 1.115969e+02
## 18 -6.284815e+00
## 19 -7.493326e+00
## 20 -9.749210e+00
## 21 -3.595066e-01
## 22 4.258792e-02
## 23 8.828906e-03
## 24 7.082474e-01
## 25 1.654037e+01
## 26 -6.174070e+00
## 27 -3.536990e-01
## 28 -2.249601e+00
## 29 7.635151e-01
## 30 -3.622119e-02
## 31 6.060651e-03
## 32 -3.340233e-01
## poly(dis, degree = 20, raw = TRUE)10
## 1 3.659894e-03

```

```

## 2      3.659894e-03
## 3      9.060469e-01
## 4      2.667325e-07
## 5      1.524689e+00
## 6      4.398600e-14
## 7      1.524689e+00
## 8      2.050475e-02
## 9      4.041889e-02
## 10     1.174167e-03
## 11     1.174167e-03
## 12     4.051252e-05
## 13     4.051252e-05
## 14     4.051252e-05
## 15     7.818264e+02
## 16     4.694151e+02
## 17     1.884383e+02
## 18     7.708887e+00
## 19     9.372613e+00
## 20     1.255610e+01
## 21     3.208788e-01
## 22     2.999058e-02
## 23     5.220046e-03
## 24     6.816146e-01
## 25     2.259120e+01
## 26     7.558104e+00
## 27     3.151244e-01
## 28     2.461661e+00
## 29     7.409644e-01
## 30     2.505229e-02
## 31     3.436628e-03
## 32     2.957080e-01
##      poly(displacement, degree = 20, raw = TRUE)11
## 1     -2.088408e-03
## 2     -2.088408e-03
## 3     -8.971514e-01
## 4       5.870615e-08
## 5       1.590375e+00
## 6     -2.030701e-15
## 7       1.590375e+00
## 8     -1.390081e-02
## 9     -2.932532e-02
## 10    -5.980021e-04
## 11    -5.980021e-04
## 12     1.473493e-05
## 13     1.473493e-05
## 14     1.473493e-05
## 15     1.522024e+03
## 16     8.683859e+02
## 17     3.181897e+02
## 18    -9.455639e+00
## 19    -1.172322e+01
## 20    -1.617113e+01
## 21    -2.864014e-01
## 22     2.111949e-02

```

```

## 23          3.086326e-03
## 24          6.559833e-01
## 25          3.085554e+01
## 26         -9.252395e+00
## 27         -2.807568e-01
## 28         -2.693712e+00
## 29          7.190798e-01
## 30         -1.732735e-02
## 31          1.948703e-03
## 32         -2.617878e-01
##   poly(dis, degree = 20, raw = TRUE)12
## 1          1.191687e-03
## 2          1.191687e-03
## 3          8.883432e-01
## 4          1.292085e-08
## 5          1.658890e+00
## 6          9.375132e-17
## 7          1.658890e+00
## 8          9.423787e-03
## 9          2.127655e-02
## 10         3.045620e-04
## 11         3.045620e-04
## 12         5.359289e-06
## 13         5.359289e-06
## 14         5.359289e-06
## 15         2.963005e+03
## 16         1.606455e+03
## 17         5.372829e+02
## 18         1.159819e+01
## 19         1.466334e+01
## 20         2.082696e+01
## 21         2.556285e-01
## 22         1.487243e-02
## 23         1.824774e-03
## 24         6.313158e-01
## 25         4.214316e+01
## 26         1.132649e+01
## 27         2.501373e-01
## 28         2.947637e+00
## 29         6.978415e-01
## 30         1.198442e-02
## 31         1.104992e-03
## 32         2.317585e-01
##   poly(dis, degree = 20, raw = TRUE)13
## 1         -6.800002e-04
## 2         -6.800002e-04
## 3         -8.796216e-01
## 4          2.843798e-09
## 5          1.730357e+00
## 6         -4.328215e-18
## 7          1.730357e+00
## 8         -6.388677e-03
## 9         -1.543689e-02
## 10        -1.551132e-04

```

```

## 11 -1.551132e-04
## 12 1.949243e-06
## 13 1.949243e-06
## 14 1.949243e-06
## 15 5.768242e+03
## 16 2.971832e+03
## 17 9.072354e+02
## 18 -1.422621e+01
## 19 -1.834083e+01
## 20 -2.682325e+01
## 21 -2.281620e-01
## 22 1.047322e-02
## 23 1.078889e-03
## 24 6.075759e-01
## 25 5.756003e+01
## 26 -1.386554e+01
## 27 -2.228572e-01
## 28 -3.225498e+00
## 29 6.772306e-01
## 30 -8.288992e-03
## 31 6.265739e-04
## 32 -2.051738e-01
## poly(displacement, degree = 20, raw = TRUE)14
## 1 3.880216e-04
## 2 3.880216e-04
## 3 8.709855e-01
## 4 6.259021e-10
## 5 1.804903e+00
## 6 1.998206e-19
## 7 1.804903e+00
## 8 4.331081e-03
## 9 1.120000e-02
## 10 7.899901e-05
## 11 7.899901e-05
## 12 7.089653e-07
## 13 7.089653e-07
## 14 7.089653e-07
## 15 1.122935e+04
## 16 5.497686e+03
## 17 1.531923e+03
## 18 1.744972e+01
## 19 2.294061e+01
## 20 3.454593e+01
## 21 2.036467e-01
## 22 7.375287e-03
## 23 6.378875e-04
## 24 5.847287e-01
## 25 7.861673e+01
## 26 1.697376e+01
## 27 1.985523e-01
## 28 3.529552e+00
## 29 6.572283e-01
## 30 5.733060e-03
## 31 3.552921e-04

```

```

## 32          1.816387e-01
## poly(dis, degree = 20, raw = TRUE)15
## 1          -2.214128e-04
## 2          -2.214128e-04
## 3          -8.624343e-01
## 4           1.377571e-10
## 5           1.882660e+00
## 6          -9.225115e-21
## 7           1.882660e+00
## 8          -2.936174e-03
## 9          -8.125996e-03
## 10         -4.023413e-05
## 11         -4.023413e-05
## 12          2.578600e-07
## 13          2.578600e-07
## 14          2.578600e-07
## 15          2.186077e+04
## 16          1.017034e+04
## 17          2.586746e+03
## 18         -2.140364e+01
## 19         -2.869400e+01
## 20         -4.449205e+01
## 21         -1.817655e-01
## 22          5.193706e-03
## 23          3.771477e-04
## 24          5.627407e-01
## 25          1.073764e+02
## 26         -2.077874e+01
## 27         -1.768980e-01
## 28         -3.862268e+00
## 29          6.378169e-01
## 30         -3.965256e-03
## 31          2.014646e-04
## 32         -1.608032e-01
## poly(dis, degree = 20, raw = TRUE)16
## 1           1.263425e-04
## 2           1.263425e-04
## 3           8.539670e-01
## 4           3.031947e-11
## 5           1.963768e+00
## 6           4.258957e-22
## 7           1.963768e+00
## 8           1.990523e-03
## 9           5.895695e-03
## 10          2.049121e-05
## 11          2.049121e-05
## 12          9.378705e-08
## 13          9.378705e-08
## 14          9.378705e-08
## 15          4.255754e+04
## 16          1.881444e+04
## 17          4.367880e+03
## 18          2.625347e+01
## 19          3.589030e+01

```

```

## 20          5.730175e+01
## 21          1.622354e-01
## 22          3.657429e-03
## 23          2.229867e-04
## 24          5.415795e-01
## 25          1.466570e+02
## 26          2.543669e+01
## 27          1.576055e-01
## 28          4.226348e+00
## 29          6.189788e-01
## 30          2.742559e-03
## 31          1.142384e-04
## 32          1.423577e-01
##      poly(dis, degree = 20, raw = TRUE)17
## 1      -7.209356e-05
## 2      -7.209356e-05
## 3      -8.455828e-01
## 4          6.673124e-12
## 5          2.048369e+00
## 6      -1.966232e-23
## 7          2.048369e+00
## 8      -1.349437e-03
## 9      -4.277534e-03
## 10     -1.043616e-05
## 11     -1.043616e-05
## 12          3.411158e-08
## 13          3.411158e-08
## 14          3.411158e-08
## 15          8.284906e+04
## 16          3.480543e+04
## 17          7.375435e+03
## 18     -3.220223e+01
## 19     -4.489141e+01
## 20     -7.379949e+01
## 21     -1.448037e-01
## 22          2.575576e-03
## 23          1.318397e-04
## 24          5.212140e-01
## 25          2.003073e+02
## 26     -3.113880e+01
## 27     -1.404169e-01
## 28     -4.624748e+00
## 29          6.006970e-01
## 30     -1.896884e-03
## 31          6.477766e-05
## 32     -1.260280e-01
##      poly(dis, degree = 20, raw = TRUE)18
## 1          4.113801e-05
## 2          4.113801e-05
## 3          8.372809e-01
## 4          1.468713e-12
## 5          2.136615e+00
## 6          9.077497e-25
## 7          2.136615e+00

```

```

## 8          9.148253e-04
## 9          3.103501e-03
## 10         5.315126e-06
## 11         5.315126e-06
## 12         1.240683e-08
## 13         1.240683e-08
## 14         1.240683e-08
## 15         1.612867e+05
## 16         6.438767e+04
## 17         1.245388e+04
## 18         3.949891e+01
## 19         5.614994e+01
## 20         9.504710e+01
## 21         1.292450e-01
## 22         1.813731e-03
## 23         7.794958e-05
## 24         5.016144e-01
## 25         2.735840e+02
## 26         3.811915e+01
## 27         1.251030e-01
## 28         5.060704e+00
## 29         5.829553e-01
## 30         1.311975e-03
## 31         3.673149e-05
## 32         1.115716e-01
##      poly(displ, degree = 20, raw = TRUE)19
## 1         -2.347417e-05
## 2         -2.347417e-05
## 3         -8.290606e-01
## 4          3.232544e-13
## 5          2.228663e+00
## 6         -4.190806e-26
## 7          2.228663e+00
## 8         -6.201884e-04
## 9         -2.251699e-03
## 10        -2.706989e-06
## 11        -2.706989e-06
## 12         4.512525e-09
## 13         4.512525e-09
## 14         4.512525e-09
## 15         3.139855e+05
## 16         1.191128e+05
## 17         2.102914e+04
## 18        -4.844893e+01
## 19        -7.023205e+01
## 20        -1.224121e+02
## 21        -1.153581e-01
## 22         1.277237e-03
## 23         4.608730e-05
## 24         4.827518e-01
## 25         3.736669e+02
## 26        -4.666427e+01
## 27        -1.114592e-01
## 28        -5.537755e+00

```



```

## [5,] 1 2.453e-03 575.00000
## [6,] 1 2.691e-03 523.90000
## [7,] 1 2.953e-03 477.40000
## [8,] 1 3.240e-03 435.00000
## [9,] 1 3.555e-03 396.30000
## [10,] 1 3.900e-03 361.10000
## [11,] 1 4.278e-03 329.00000
## [12,] 1 4.694e-03 299.80000
## [13,] 1 5.149e-03 273.20000
## [14,] 1 5.648e-03 248.90000
## [15,] 1 6.195e-03 226.80000
## [16,] 1 6.795e-03 206.60000
## [17,] 1 7.452e-03 188.30000
## [18,] 1 8.172e-03 171.60000
## [19,] 1 8.962e-03 156.30000
## [20,] 1 9.826e-03 142.40000
## [21,] 1 1.077e-02 129.80000
## [22,] 1 1.181e-02 118.20000
## [23,] 1 1.295e-02 107.70000
## [24,] 1 1.419e-02 98.17000
## [25,] 1 1.555e-02 89.45000
## [26,] 1 1.704e-02 81.50000
## [27,] 1 1.867e-02 74.26000
## [28,] 1 2.045e-02 67.66000
## [29,] 1 2.240e-02 61.65000
## [30,] 1 2.453e-02 56.18000
## [31,] 1 2.685e-02 51.19000
## [32,] 1 2.939e-02 46.64000
## [33,] 1 3.215e-02 42.50000
## [34,] 1 3.517e-02 38.72000
## [35,] 1 3.846e-02 35.28000
## [36,] 1 4.205e-02 32.15000
## [37,] 1 4.595e-02 29.29000
## [38,] 1 5.019e-02 26.69000
## [39,] 1 5.480e-02 24.32000
## [40,] 1 5.981e-02 22.16000
## [41,] 1 6.523e-02 20.19000
## [42,] 1 7.111e-02 18.40000
## [43,] 1 7.748e-02 16.76000
## [44,] 1 8.435e-02 15.27000
## [45,] 1 9.177e-02 13.92000
## [46,] 1 9.976e-02 12.68000
## [47,] 1 1.084e-01 11.55000
## [48,] 1 1.176e-01 10.53000
## [49,] 1 1.275e-01 9.59100
## [50,] 1 1.381e-01 8.73900
## [51,] 1 1.494e-01 7.96300
## [52,] 1 1.614e-01 7.25500
## [53,] 1 1.741e-01 6.61100
## [54,] 1 1.876e-01 6.02400
## [55,] 1 2.019e-01 5.48900
## [56,] 1 2.168e-01 5.00100
## [57,] 1 2.325e-01 4.55700
## [58,] 1 2.489e-01 4.15200

```

```
## [59,] 1 2.660e-01 3.78300
## [60,] 1 2.837e-01 3.44700
## [61,] 1 3.019e-01 3.14100
## [62,] 1 3.206e-01 2.86200
## [63,] 1 3.398e-01 2.60700
## [64,] 1 3.592e-01 2.37600
## [65,] 1 3.789e-01 2.16500
## [66,] 1 3.987e-01 1.97200
## [67,] 1 4.184e-01 1.79700
## [68,] 1 4.381e-01 1.63800
## [69,] 1 4.575e-01 1.49200
## [70,] 1 4.766e-01 1.36000
## [71,] 1 4.953e-01 1.23900
## [72,] 1 5.134e-01 1.12900
## [73,] 1 5.308e-01 1.02800
## [74,] 1 5.475e-01 0.93710
## [75,] 1 5.633e-01 0.85380
## [76,] 1 5.783e-01 0.77800
## [77,] 1 5.924e-01 0.70890
## [78,] 1 6.056e-01 0.64590
## [79,] 1 6.178e-01 0.58850
## [80,] 1 6.290e-01 0.53620
## [81,] 1 6.393e-01 0.48860
## [82,] 1 6.487e-01 0.44520
## [83,] 1 6.572e-01 0.40560
## [84,] 1 6.648e-01 0.36960
## [85,] 1 6.717e-01 0.33680
## [86,] 1 6.778e-01 0.30690
## [87,] 1 6.832e-01 0.27960
## [88,] 1 6.880e-01 0.25480
## [89,] 1 6.922e-01 0.23210
## [90,] 1 6.959e-01 0.21150
## [91,] 1 6.991e-01 0.19270
## [92,] 1 7.019e-01 0.17560
## [93,] 1 7.043e-01 0.16000
## [94,] 1 7.064e-01 0.14580
## [95,] 1 7.082e-01 0.13280
## [96,] 1 7.097e-01 0.12100
## [97,] 1 7.111e-01 0.11030
## [98,] 1 7.122e-01 0.10050
## [99,] 1 7.131e-01 0.09155
## [100,] 1 7.140e-01 0.08342
```

```
coef_pen_mod1 <- tibble(
  lambda = penalized_mod1$lambda,
  intercept = coef(penalized_mod1)[1, ],
  disp = coef(penalized_mod1)[3, ]
)
coef_pen_mod1 # Gives estimated coefficients for each value of lambda
```

```
## # A tibble: 100 x 3
##   lambda intercept      disp
##   <dbl>      <dbl>    <dbl>
## 1  834.  4.60e-17 -8.56e-37
## 2  760.  4.59e-17 -1.10e- 3
```

```
## 3 693. 4.59e-17 -1.20e- 3
## 4 631. 4.59e-17 -1.32e- 3
## 5 575. 4.58e-17 -1.45e- 3
## 6 524. 4.58e-17 -1.59e- 3
## 7 477. 4.58e-17 -1.74e- 3
## 8 435. 4.58e-17 -1.91e- 3
## 9 396. 4.58e-17 -2.10e- 3
## 10 361. 4.58e-17 -2.30e- 3
## # ... with 90 more rows
```


```
# Higher lambda ==> stronger penalty ==> smaller estimates.
```

```
View(coef_pen_mod1)
```

```
# As lambda gets smaller, the estimated coefficient approaches  
# what it was in the unpenalized model.
```

```
# Plot the disp coefficient as a function of lambda and compare to  
# our unpenalized model
```

```
coef_pen_mod1 %>%  
  ggplot(aes(x = lambda, y = disp)) +  
  theme_light() +  
  geom_line() +  
  geom_hline(yintercept = coef(mod1)[2], linetype = "dashed", colour = "red") +  
  coord_cartesian(xlim = c(0,10))
```



```
# Low lambda ==> limited penalization ==> answer is close to unpenalized
```

```
# High lambda ==> a lot of penalization ==> answer doesn't want to move far from zero
```

```
# Try the quadratic one yourself.
```

```
# Try the big one:
```

```
penalized_mod3 <- glmnet(x = X3,y = y,alpha = 0,nlambda = 100)
coef_pen_mod3 <- tibble(lambda = penalized_mod3$lambda) %>%
  bind_cols(as_tibble(as.matrix(t(coef(penalized_mod3)))))
```

```
## Warning: `as_tibble.matrix()` requires a matrix with column names or a `.name_repair` argument. Using
## This warning is displayed once per session.
```

```
options(scipen = 999) # Turn off scientific notation
View(coef_pen_mod3) # Opens in a spreadsheet
```

```
# Pick any lambda- I chose the smallest one. You can try others!
```

```
whichlambda <- penalized_mod3$lambda[length(penalized_mod3$lambda)]
```

```
# Plot predictions
```

```
mtscaled %>%
```

```
  mutate(predvals = predict(penalized_mod3,newx = X3,s = whichlambda)) %>%
```

```
  ggplot(aes(x=dis,y=mpg)) +
```


```
  theme_light() +
```

```
  geom_point(colour="black",fill="red",pch=21,size=2) +
```

```
  geom_line(aes(y = predvals),colour = "purple") +
```

```
  labs(subtitle="Linear Regression Model with 20-degree polynomial terms fit by penalized regression")
```

Linear Regression Model with 20-degree polynomial terms fit by penalized regression



```
# Looks a lot like the quadratic, except it's fine at the
# right end point.
# Very little effort on our part to build this very
```

```

# sophisticated model.
# The betas here should have higher variance than if we had
# just fit a linear/quadratic model. We pay a price for
# generality. However, they have MUCH lower variance than
# the original 20-degree polynomial we tried.
# "No such thing as a free lunch". We pay a small price for
# being able to let the data tell us what model to use.
# I think it's worth it...

# Advanced: plot it for ALL the lambda values together! Woah!
# Requires the purrr package
# THIS WILL NOT BE TESTED
library(purrr)

##
## Attaching package: 'purrr'

## The following objects are masked from 'package:foreach':
##
##      accumulate, when

predframe <- purrr::map(penalized_mod3$lambda,
  ~mtscaled %>%
    mutate(predvals = predict(penalized_mod3,newx = X3,s = .x)[,1],
      lambda = .x)) %>%
  purrr::reduce(rbind)


predframe

## # A tibble: 3,200 x 4
##       disp    mpg predvals lambda
##       <dbl> <dbl>   <dbl>   <dbl>
## 1 -0.571  0.151 4.60e-17  834.
## 2 -0.571  0.151 4.60e-17  834.
## 3 -0.990  0.450 4.60e-17  834.
## 4  0.220  0.217 4.60e-17  834.
## 5  1.04   -0.231 4.60e-17  834.
## 6 -0.0462 -0.330 4.60e-17  834.
## 7  1.04   -0.961 4.60e-17  834.
## 8 -0.678  0.715 4.60e-17  834.
## 9 -0.726  0.450 4.60e-17  834.
## 10 -0.509 -0.148 4.60e-17  834.
## # ... with 3,190 more rows

predframe %>%
  ggplot(aes(x=disp,y=mpg,group = lambda)) +
  theme_light() +
  geom_point(colour="black",fill="red",pch=21,size=2) +
  geom_line(aes(y = predvals),colour = "grey",alpha = .1) +
  labs(subtitle="Linear Regression Model with 20-degree polynomial terms fit by penalized regression")

```

Linear Regression Model with 20-degree polynomial terms fit by penalized regression



*# That's a comparison of all the curves fit using different
values of lambda.*